

Image Compression (1)

Qiufeng Wang

Qiufeng.Wang@xjtlu.edu.cn

Outline

- Introduction to compression (source coding)
Why, what, how compression
- Concept of lossless and lossy compression
- Types of redundancy
- Fixed-length codes
- Variable length codes

Why image compression?

---Data storage

- Comparison of common broadcast resolutions

Format ↕	Resolution ↕	Display aspect ratio ↕	Pixels ↕
Ultra-high-definition television	3840 × 2160	1.78:1 (16:9)	8,294,400
Ultra-wide-television	5120 × 2160	2.37:1 (21:9)	11,059,200
DCI 4K (native resolution)	4096 × 2160	1.90:1 (256:135)	8,847,360
DCI 4K (CinemaScope cropped)	4096 × 1716	2.39:1 (1024:429)	7,028,736
DCI 4K (flat cropped)	3996 × 2160	1.85:1 (999:540)	8,631,360

Why image compression?

---Data storage

- Comparison of common broadcast resolutions

Format	Resolution	Display aspect ratio	Pixels
Ultra-high-definition television	3840 × 2160	1.78:1 (16:9)	8,294,400
Ultra-wide-television	5120 × 2160	2.37:1 (21:9)	11,059,200
DCI 4K (native resolution)	4096 × 2160	1.90:1 (256:135)	8,847,360
DCI 4K (CinemaScope cropped)	4096 × 1716	2.39:1 (1024:429)	7,028,736
DCI 4K (flat cropped)	3996 × 2160	1.85:1 (999:540)	8,631,360

One image : $4096 \times 2160 \times 3 = 26.5$ M Byte

One hour video : $4096 \times 2160 \times 3 \times 30 \times 3600 = 2862$ GByte

Why image compression?

---Data storage

- Image data can take **a lot of space**. An **example** of storage requirement for a full-frame camera (Nikon D600 - RAW 14 bit):

$$\begin{array}{lll} (6016 \times 4016) & \times 14 & = 338243584 \text{ bits} \\ \text{Spatial resolution} & \text{bit/channel} & \\ & & = 42\text{MB} \end{array}$$

Why image compression?

---Data usage pattern

- Symmetric applications
 - Data is typically compressed and decompressed the same **number of times**
 - Compression and decompression involves roughly the **same complexity**
 - Example: video telephony



Why image compression?

---Data usage pattern

- Asymmetric applications
 - Data is typically compressed once and **decompressed many times**
 - Compression is computationally expensive, **decompression is simple**
 - Movie compression

Why Image Compression

---Benefits

- Image **compression** has the following **benefits**:
 - Less memory for **storage**.
 - Reduced **transmission time** and **costs**.
 - Reduced the required **memory** for computation.

Why image compression?

--- Redundancy



BMP: 258KB



JPG: 82KB

What is Compression

- The **fundamental task** of image compression is to **reduce** the amount of data (**bytes**) required to **represent** an image (**information**).
 - We do this by **removing** image data **redundancies**.


Data vs. information

- Data (bytes) is not equal to information
- Data is the means by which information is conveyed
 - The same story can be told with a different number of words if the teller is long-winded or short and to the point!
- There are different ways to represent the same information.

Data vs. information

- Information theory address one essential problem:
How much data is needed in order to store a certain amount of information?

Claude Shannon



Born	April 30, 1916 Petoskey, Michigan, United States
Died	February 24, 2001 (aged 84) Medford, Massachusetts, United States
Nationality	American

Measuring information

- A discrete *memoryless* source generates symbols from a set X of M elements (alphabet); each symbol is characterized by its probability of occurrence p_i

$$X = \{x_i\}_{i=1}^M$$

$$\{p_i\}_{i=1}^M$$

- How do we measure the amount of information carried by message x_i ?

Properties of information

- The amount of information carried by a message is inversely proportional to its probability

$$I(x_j) > I(x_i) \quad \text{if} \quad p_j < p_i$$

- Statistically independent messages:

$$P(x_i, x_j) = P(x_i)P(x_j) \Rightarrow I(x_i, x_j) = I(x_i) + I(x_j)$$

Definition of “information”

$$I(x_i) = \log_2 \frac{1}{p_i}$$

and is measured in bits

- Average amount of information carried by the memoryless source: *first order entropy* (bits/symbol)

$$H(X) = \sum_{i=1}^M p_i I(x_i) = \sum_{i=1}^M p_i \log_2 \frac{1}{p_i}$$

Examples

$$X = \begin{array}{ll} x_1 & p_1 = 1/2 \\ x_2 & p_2 = 1/4 \\ x_3 & p_3 = 1/8 \\ x_4 & p_4 = 1/8 \end{array}$$

$$H(X) = 1.75 \text{ bits/symbol}$$

$$X = \begin{array}{ll} x_1 & p_1 = 1/4 \\ x_2 & p_2 = 1/4 \\ x_3 & p_3 = 1/4 \\ x_4 & p_4 = 1/4 \end{array}$$

$$H(X) = 2 \text{ bits/symbol}$$

Equiprobable symbols carry more information, and
are more difficult to compress

Bounds on Entropy

- **Theorem:** the first order entropy of a memoryless M-symbol alphabet is limited by

$$H(X) \leq \log_2 M$$

- Example: 8 bit quantizer ($M=2^8$)

$$H(X) \leq 8 \text{ bit/symbol}$$

with the equality if the symbols are equiprobable

Entropy ...

The first order
Entropy of the
source (H)

- What is the name of this value:

$$-\sum_{i=1}^M p_i \log_2 p_i = \sum_{i=1}^M p_i \log_2 \frac{1}{p_i} [\text{bit} / \text{symbol}]$$

- The average amount of information carried by a source is the *entropy*
- The amount of information carried by a symbol is $\log_2 \frac{1}{p_i}$
- $\log_2 \frac{1}{p_i}$ is the **uncertainty** in symbol e_i (or the “surprise” when we see this symbol). **Entropy** – average “surprise”.

Entropy of images

- Entropy in Matlab :

» `h = entropy(uint8(im))`

1st order Entropy = 7.4



1st order Entropy = 6.2



White is the most likely value in this picture

How to quantify compression efficiency ?

- Since there are different ways to represent the same information, then **how to quantify** which **representation** is **better** in term of compression efficiency ?

Compression ratio

Size (bmp) = 17462




Size (bmp) = 17462




- The monochrome image is 128x128 pixels (**symbols of information**)

Size (bmp) = 17462



Size (bmp) = 17462



Compression ratio

Size (jpg) = 10872,




Size (jpg) = 6972,




- The monochrome image is 128x128 pixels (**symbols of information**)

Size (jpg) = 4508,



Size (jpg) = 2762,



Compression ratio

Size (jpg) = 10872, CR = 1.6




Size (jpg) = 6972, CR = 2.5




6.3 = 17462/2762

Size (jpg) = 4508, CR = 3.9



Size (jpg) = 2762, CR = 6.3



- The monochrome image is 128x128 pixels (**symbols of information**)
- Let n_1 and n_2 be the number of **data-unit** (**coding words**) in two data sets that represents the same information then

$$CR = \frac{n_1}{n_2}$$

Compression ratio

- A **compression** algorithm **searches** a **representation** with:

$$CR > 1$$

Compression types

- There are **two** major **categories** of compression **algorithms**

1. **Lossless** compression algorithms:

- The **original data** is recovered **perfectly** after decompression.
- There is a **theoretical limit** on maximal compression.
Depend purely on the **image content**.
- **Practical** compression ratios **CR<10** (**still images**).
- Where it is used ? Medical, Hyperspectral, aerial, ...

Compression types

- There are **two** major **categories** of compression **algorithms**

2. **Lossy** compression algorithms:

- **Decompression** results in an **approximation** of the original image.
- **Compression rate** is a **function** of reconstruction **quality**.
- Practical compression ratios **CR >10** (still images).
- Where it is used ? Application where the end user is human who will view the media content.

Lossless - Introduction

- The goal of lossless compression is

to minimize the average length
of the compressed symbols

exploiting statistical properties of the data

- *probability distribution*
- *correlation (redundancy) of the data*

How to compress data?

- How to compress data?
- This could be achieved by finding a more compact signal representation by reducing the intrinsic redundancy (duplication) of the source signal

Types of redundancy

Types of redundancy

- Typical signals contain redundancy which could be exploited:
 - Coding (data representation) redundancy
 - Interpixel redundancy (Correlation between adjacent samples)
 - Psychovisual redundancy.


Types of redundancy

- In general, **coding** and **interpixel redundancies** can be exploited **without losing** any information → lossless compression.
- In general, **lossy** compression achieves **much higher compression rates** than lossless compression by **exploiting** the **psychovisual redundancy**

Coding redundancy

- This type of redundancy deals with the way we represent data.

How to represent the information ?



8 bit/symbol (pixel)

- How to write this image on a file ?
- How are we going to represent each pixel ?

Fixed-Length Code

0000 0000	0000 0001	0000 0000	0000 0000
0000 0000	0000 0011	0000 0000	0000 0011
0000 0010	0000 0000	0000 0011	0001
0000 0000	0000 0011	0000 0010	0000 0000

8 bit/symbol

- Is it efficient to use 8 bits per symbol for such image ?

Fixed-Length Code

0000 0000	0000 0001	0000 0000	0000 0000
0000 0000	0000 0011	0000 0000	0000 0011
0000 0010	0000 0000	0000 0011	0001
0000 0000	0000 0011	0000 0010	0000 0000

8 bit/symbol

00	01	00	00
00	11	00	11
10	00	11	01
00	11	10	00

2 bit/symbol

- Which one is more **efficient** (compact) **representation** of the information ?

Fixed-Length Codes

- Properties
 - Use the **same number** of bits to represent all **possible** symbols (pixels) produced by the **source**
 - **Simplify** the **decoding** process
- Examples
 - American Standard Code for Information Interchange (ASCII) code
 - Bar codes
 - Universal Product Code (UPC) on products in stores
 - Credit card codes



ASCII Code

- ASCII is used to **encode** and communicate **alphanumeric** characters for plain text; 7 bits per character

Dec	Hx	Oct	Char	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr	Dec	Hx	Oct	Html	Chr
0	0	000	NUL (null)	32	20	040	Space	64	40	100	@	96	60	140	`			
1	1	001	SOH (start of heading)	33	21	041	!	65	41	101	A	97	61	141	a			
2	2	002	STX (start of text)	34	22	042	"	66	42	102	B	98	62	142	b			
3	3	003	ETX (end of text)	35	23	043	#	67	43	103	C	99	63	143	c			
4	4	004	EOT (end of transmission)	36	24	044	\$	68	44	104	D	100	64	144	d			
5	5	005	ENQ (enquiry)	37	25	045	%	69	45	105	E	101	65	145	e			
6	6	006	ACK (acknowledge)	38	26	046	&	70	46	106	F	102	66	146	f			
7	7	007	BEL (bell)	39	27	047	'	71	47	107	G	103	67	147	g			
8	8	010	BS (backspace)	40	28	050	(72	48	110	H	104	68	150	h			
9	9	011	TAB (horizontal tab)	41	29	051)	73	49	111	I	105	69	151	i			
10	A	012	LF (NL line feed, new line)	42	2A	052	*	74	4A	112	J	106	6A	152	j			
11	B	013	VT (vertical tab)	43	2B	053	+	75	4B	113	K	107	6B	153	k			
12	C	014	FF (NP form feed, new page)	44	2C	054	,	76	4C	114	L	108	6C	154	l			
13	D	015	CR (carriage return)	45	2D	055	-	77	4D	115	M	109	6D	155	m			
14	E	016	SO (shift out)	46	2E	056	.	78	4E	116	N	110	6E	156	n			
15	F	017	SI (shift in)	47	2F	057	/	79	4F	117	O	111	6F	157	o			
16	10	020	DLE (data link escape)	48	30	060	0	80	50	120	P	112	70	160	p			
17	11	021	DC1 (device control 1)	49	31	061	1	81	51	121	Q	113	71	161	q			
18	12	022	DC2 (device control 2)	50	32	062	2	82	52	122	R	114	72	162	r			
19	13	023	DC3 (device control 3)	51	33	063	3	83	53	123	S	115	73	163	s			
20	14	024	DC4 (device control 4)	52	34	064	4	84	54	124	T	116	74	164	t			
21	15	025	NAK (negative acknowledge)	53	35	065	5	85	55	125	U	117	75	165	u			
22	16	026	SYN (synchronous idle)	54	36	066	6	86	56	126	V	118	76	166	v			
23	17	027	ETB (end of trans. block)	55	37	067	7	87	57	127	W	119	77	167	w			
24	18	030	CAN (cancel)	56	38	070	8	88	58	130	X	120	78	170	x			
25	19	031	EM (end of medium)	57	39	071	9	89	59	131	Y	121	79	171	y			
26	1A	032	SUB (substitute)	58	3A	072	:	90	5A	132	Z	122	7A	172	z			
27	1B	033	ESC (escape)	59	3B	073	;	91	5B	133	[123	7B	173	{			
28	1C	034	FS (file separator)	60	3C	074	<	92	5C	134	\	124	7C	174				
29	1D	035	GS (group separator)	61	3D	075	=	93	5D	135]	125	7D	175	}			
30	1E	036	RS (record separator)	62	3E	076	>	94	5E	136	^	126	7E	176	~			
31	1F	037	US (unit separator)	63	3F	077	?	95	5F	137	_	127	7F	177	DEL			

Source: www.asciitable.com

Fixed-Length Code

00	01	00	00
00	11	00	11
10	00	11	01
00	11	10	00

- 4 different symbols we use 2 bit/symbol (fixed-length code)
- Is it the most efficient representation of this information ?

.....-Length Codes

00	01	00	00
00	11	00	11
10	00	11	01
00	11	10	00

2 bit/symbol

1	001	1	1
1	01	1	01
000	1	01	001
1	01	000	1


(28/16) bit/symbol

- Which one is the most efficient code ?
- Main **problem** with **fixed-length codes**:
inefficiency

Variable-Length Codes


- Main **properties** of variable-length codes (**VLC**)
 - Use a **different number** of bits to **represent** each **symbol**
 - Allocate **shorter-length codewords** to symbols that occur **more frequently**
 - Allocate **longer-length codewords** to **rarely-occurred** symbols
 - More **efficient representation**; good for compression

Variable-Length Codes







- Encoder:
 - Associate the **shortest** binary codeword to the most “probable” **symbols**

Variable-Length Codes



Symbols

	1
	01
	000
	001

Dictionary

1	001	1	1
1	01	1	01
000	1	01	001
1	01	000	1

encode





- Encoder:
 - Associate the **shortest** binary codeword to the most “probable” **symbols**
 - Generate a dictionary, encode and send all

Decoding VLC

- Suppose you **received** a file with these data (and **with** the dictionary):

1100010010110111010001010011


- To **generate** the **information** (decode)...

	1
	01
	000
	001

Decoding VLC

- Suppose you **received** a file with these data (and **with** the dictionary):

1 1 0 0 0 1 0 0 1 0 1 1 0 1 1 1 0 1 0 0 0 1 0 1 0 0 1 1




- To generate the information (decode)...

1 1 000 1 001 01 1 01 1 1 01 000 1 01 001 1


Decoding VLC

- Suppose you **received** a file with these data (and **with** the dictionary):

1 1 0 0 0 1 0 0 1 0 1 1 0 1 1 0 1 0 0 0 1 0 1 0 0 1 1




- To generate the information (decode)...



The **4x4 image** write **column-wise**

Try decoding this string

- try it ..



1 1 0 0 0 1 0 0 1 1 1 1 1 1 1 1 1 0 0 0 1 1 1 0 0 1 1

Variable length coding

- Associate the **shortest** binary **codewords** to the most “probable” **symbols**

■	1	8/16
■	01	4/16
■	000	2/16
■	001	2/16

- Prefix rule:** no codeword should be a prefix of a longer codeword, e.g., 1 is not prefix of 01

1 1 0 0 0 1 0 0 1 0 1 1 0 1 1 0 1 0 0 0 1 0 1 0 0 1 1



1 1 000 1 001 01 1 01 1 1 01 000 1 01 001 1

Variable-Length Codes

- Examples of VLC
 - Shannon-Fano code
 - Huffman code

Shannon-Fano Code

- Algorithm

- Line up symbols by decreasing probability of occurrence
- Divide symbols into 2 groups so that both have similar combined probability
- Assign **0** to 1st group and **1** to the 2nd
- Repeat step 2

- Example


Symbols	Prob.	Code-word
A	0.35	00
B	0.17	01
C	0.17	10
D	0.16	110
E	0.15	111

$$\begin{aligned}\text{Average code-word length} &= \\ &0.35 \times 2 + 0.17 \times 2 + 0.17 \times 2 \\ &\quad + 0.16 \times 3 + 0.15 \times 3 \\ &= 2.31 \text{ bits per symbol}\end{aligned}$$

Huffman Coding

Huffman Coding

- Allows us to build a **quasi-optimal** VLC



Huffman Coding

- Building a Huffman tree
 - Order the symbols by the **descending** order of probabilities



a 0.5

b 0.25

c 0.125

d 0.125

Huffman Coding

- Building a Huffman tree

- Then **combine** the symbols of **lowest probability** to form new symbols **recursively**



a 0.5

b 0.25

c 0.125


d 0.125

0.25




Huffman Coding

- Building a Huffman tree




Huffman Coding

- Building a Huffman tree



Huffman Coding


- Code assignment



- When Huffman tree is **completely** built, then **assign** a **code-bit** to the symbols in **each branch** of the tree


Huffman Coding

- Code assignment




Huffman Coding

- Code assignment



Huffman Coding

- Allows us to construct a VLC



Huffman Coding

- In the **previous example** the average Length of the VLC code is:

$$0.5 \times 1 + 0.25 \times 2 + 0.125 \times 3 + 0.125 \times 3 = \underline{1.75} \text{ [bit / symbol]}$$

- Compression Ratio: $CR = 2 / 1.75 \approx \underline{1.14}$

Example II: Huffman Coding

- Building a Huffman tree

a 0.5

b 0.2

c 0.2

d 0.1


Example II: Huffman Coding

- Building a Huffman tree




Example II: Huffman Coding

- Building a Huffman tree



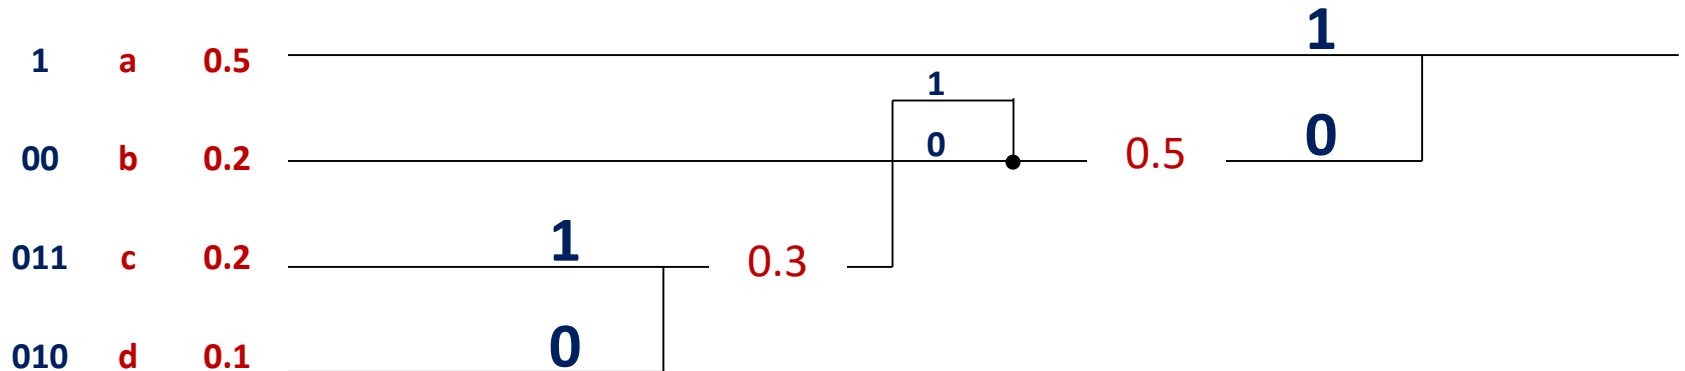
Example II: Huffman Coding

- Building a Huffman tree



Example II: Huffman Coding

- Code assignment




Huffman Code

- Shannon-Fano code [1949]
 - Top-down algorithm: assigning code from most frequent to least frequent
 - VLC, uniquely & instantaneously decodable (no code-word is a prefix of another)
 - Unfortunately not optimal in term of minimum redundancy
- Huffman code [1952]
 - Quite similar to Shannon-Fano in VLC concept
 - Bottom-up algorithm: assigning code from least frequent to most frequent
 - Minimum redundancy when probabilities of occurrence are powers-of-two
 - In JPEG images, DVD movies, MP3 music

Huffman Coding Algorithm

- Encoding algorithm
 - Order the symbols by decreasing probabilities
 - Starting from the bottom, assign **0** to the least probable symbol and **1** to the next least probable
 - Combine the two least probable symbols into one composite symbol
 - Reorder the list with the composite symbol
 - Repeat Step 2 until only two symbols remain in the list
- Huffman tree
 - Nodes: symbols or composite symbols
 - Branches: from each node, 0 defines one branch while 1 defines the other
- Decoding algorithm
 - Start at the root, follow the branches based on the bits received
 - When a leaf is reached, a symbol has just been decoded



Huffman Coding Example

Symbols	Prob.	
A	0.35	
B	0.17	
C	0.17	
D	0.16	1
E	0.15	0



Symbols	Prob.	
A	0.35	
DE	0.31	
B	0.17	1
C	0.17	0




Symbols	Prob.	
A	0.35	
BC	0.34	1
DE	0.31	0



Huffman Codes

A	0
B	111
C	110
D	101
E	100



Symbols	Prob.	
BCDE	0.65	1
A	0.35	0


Average code-word length =
 $0.35 \times 1 + 0.65 \times 3 = 2.30$ bits per symbol

Huffman coding example

- Image $\begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 1 & 1 & 0 & 0 \\ 0 & 7 & 1 & 1 & 1 & 1 & 0 & 0 \\ 0 & 7 & 5 & 5 & 5 & 5 & 2 & 2 \\ 0 & 7 & 0 & 0 & 0 & 0 & 2 & 2 \\ 0 & 7 & 2 & 2 & 2 & 2 & 2 & 2 \\ 0 & 0 & 2 & 2 & 2 & 2 & 2 & 2 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$

- How to do Huffman coding?

Huffman coding example



Some observations about Huffman Coding

Entropy ...

- In the first example of Huffman coding: does the equation of the average Length remind you of something you should know ?

$$\begin{aligned} &= 0.5 \times 1 \quad + 0.25 \times 2 \quad + 0.125 \times 3 \quad + 0.125 \times 3 \\ &= 0.5 \times (-\log_2 0.5) \quad + 0.25 \times (-\log_2 0.25) \quad + 0.125 \times (-\log_2 0.125) \quad + 0.125 \times (-\log_2 0.125) \\ &= -\sum_{i=1}^M p_i \log_2 p_i \end{aligned}$$

Entropy ...

- Entropy in Matlab :
» `h = entropy(uint8(im))`

1st order Entropy = 7.4



1st order Entropy = 6.2



White is the most likely value in this picture

Efficiency of Huffman Coding

- Huffman is a code which is **optimal** when all **symbols probabilities** are **integral** power of $\frac{1}{2}$

– In example I:

$$\text{average length} = 0.5 \times 1 + 0.25 \times 2 + 0.125 \times 3 + 0.125 \times 3 = 1.75$$

$$\text{entropy} = -(0.5 \times \log_2 0.5 + 0.25 \times \log_2 0.25 + 0.125 \times \log_2 0.125 + 0.125 \times \log_2 0.125) = 1.75$$

– In example II:

$$\text{average length} = 0.5 \times 1 + 0.20 \times 2 + 0.200 \times 3 + 0.100 \times 3 = 1.80$$

$$\text{entropy} = -(0.5 \times \log_2 0.5 + 0.2 \times \log_2 0.2 + 0.2 \times \log_2 0.2 + 0.1 \times \log_2 0.1) = 1.76$$

- What** if the probabilities are not **integral** power of $\frac{1}{2}$?

Bounds on lossless coding

- The average codeword length of a lossless coder cannot be less than entropy
- Entropy represents the target average number of bits/symbol of a lossless encoder
- *Coding Efficiency:*

$$\eta = H(X) / n$$

where n is the average codeword length

Context adaptive coding

- Example of context adaptive coding:
 - The **Run-Length Coding** reduces the length of a repeating character sequence

a c d e e e e e e e e u h r r r r g e

a c d #e8 u h #r4 g e

- The **Dictionary Coding** compress data by searching for repeating sequence of characters

- He quietly quit the theatre -

Self-Reading

Arithmetic coding

- Frequently used characters will be stored with fewer bits and not-so-frequently occurring characters will be stored with more bits, resulting in fewer bits used in total.
- Arithmetic coding encodes the entire message into a **single number**.
- Example:
 - ABBCAB
 - Equal probabilities
 - Fix-length: A=00, B=01 and C=10, 000101100001, ->12 bits
 - Arithmetic : ABBCAB=0.0112013->0.0010110010₂, -> 10 bits

Arithmetic coding

- Encode
 - set Low to 0
 - set High to 1
 - while there are input symbols do
 - take a symbol
 - $\text{CodeRange} = \text{High} - \text{Low}$
 - $\text{High} = \text{Low} + \text{CodeRange} * \text{HighRange}(\text{symbol})$
 - $\text{Low} = \text{Low} + \text{CodeRange} * \text{LowRange}(\text{symbol})$
 - end of while
 - output Low

- Example:
 - C A D A C D B
 - 0.5143876

A	B	C	D
0.1	0.4	0.2	0.3
[0, 0.1)	[0.1, 0.5)	[0.5, 0.7)	[0.7, 1]

1	C	[0.5, 0.7]
2	A	[0.5, 0.52]
3	D	[0.514, 0.52]
4	A	[0.514, 0.5146]
5	C	[0.5143, 0.51442]
6	D	[0.514384, 0.51442]
7	B	[0.5143876, 0.514402]

Arithmetic coding

- Decode
 - get encoded number
 - do
 - find symbol whose range straddles the encoded number
 - output the symbol
 - $\text{range} = \text{symbol.HighValue} - \text{symbol.LowValue}$
 - subtract symbol.LowValue from encoded number
 - divide encoded number by range
 - until no more symbols
- Example
 - Input: 0.5143876
 - Output: C A D A C D B

A	B	C	D
0.1	0.4	0.2	0.3
[0, 0.1)	[0.1, 0.5)	[0.5, 0.7)	[0.7, 1]

Index	encodedNumber	Range	Symbol	LowValue	HighValue
1	0.5143877	[0.5,0.7)	C	0.5	0.7
2	0.0719385	[0.0,0.1)	A	0	0.1
3	0.719385	[0.7,1]	D	0.7	1
4	0.064616667	[0.0,0.1)	A	0	0.1
5	0.646166667	[0.5,0.7)	C	0.5	0.7
6	0.730833333	[0.7,1]	D	0.7	1
7	0.102777778		B	0.1	0.5

End of Self-Reading

Types of redundancy

- Typical signals contain redundancy which could be exploited:
 - Coding (data representation) redundancy
 - Interpixel redundancy (Correlation between adjacent samples)
 - Psychovisual redundancy.


How could we exploit the context-data
for image compression ?

How could we exploit the context-data for image compression ?

- In natural images do we have **constant patterns** with **exactly** the same intensities?

How could we exploit the context-data for image compression ?

- Which one is **natural image** and which one is **synthetic** (graphical image) ?




How could we recognize that an image is natural ?

- In natural images we do **NOT** have **constant patterns** with **exactly** the same **intensities**.
- Even with the **smoothest surface** and **omnidirectional** light the **apparently-smooth** areas will have some **small intensities differences**.
 - This is why **run-length** and **dictionary coding** do **not** perform **well** for natural image **compression**.

Smoothness of images and interpixel redundancy

Interpixel redundancy

- Example



Interpixel redundancy

- Example




Images are smooth

165	172	181	186	190	196	195	201
169	176	184	187	192	193	194	195
169	173	182	187	190	193	189	190
173	177	182	185	191	189	189	188
168	173	179	182	189	187	188	190
169	170	175	180	183	184	185	189
166	169	173	176	181	180	186	184
171	168	167	176	176	180	177	181

Interpixel redundancy


- Example



Interpixel redundancy

Images have gradual change of intensity

- Example



96	99	99	95	86	95	107	121
97	96	93	107	99	98	109	124
98	96	95	98	98	95	109	118
98	99	98	97	96	95	108	125
98	97	97	95	97	97	106	121
99	100	97	99	95	101	111	125
98	104	98	99	97	99	110	125
103	103	95	96	92	107	111	128

Interpixel redundancy

- If a pixel value can be reasonably predicted from its neighboring pixels the image is said to contain interpixel redundancy.
- Interpixel redundancy depends on the resolution of the image.
 - The higher the spatial resolution of an image, the more probable it is that two neighboring pixels will belong to the same object.

Interpixel redundancy



- Large areas of the image are uniform → images are smooth with gradual change of intensity
- This means adjacent pixels are almost the same → **high correlation** among pixels
- This is **not exploited** using **variable length coding** (which works on each single pixel)

Interpixel redundancy

- How to **exploit** the **smoothness property** of images for **compression** ?



Predictive coding

Predictive coding

- Given that neighboring pixels in natural images are correlated, so we could exploit some prediction approaches to reduce the dynamics of the data.

Predictive coding

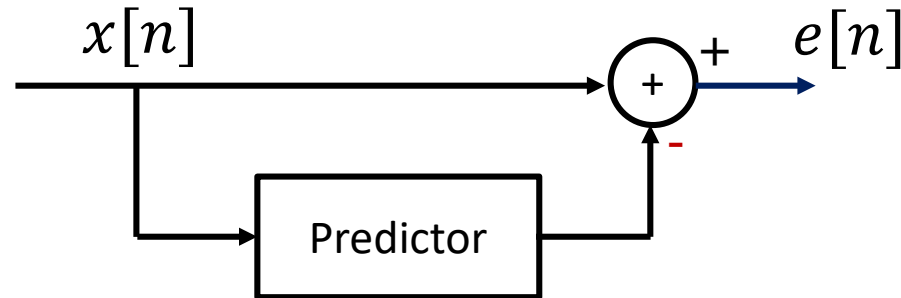
- Provides a **data representation** of the information where **code words** express the source symbol **deviations** from **predicted values** (usually values of neighboring pixels).
- Predictive coding efficiently reduces interpixel redundancies.
 - 1D & 2D – pixels are predicted from neighboring pixels.
 - Video – pixels are predicted between frames as well.

Predictive coding

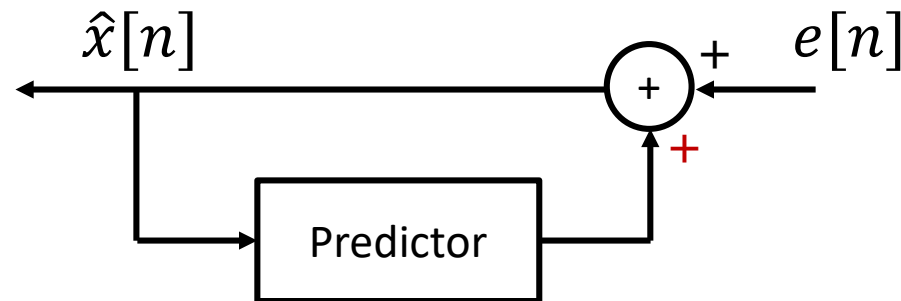
- Works well for all images with a high degree of interpixel redundancies.
- Works in the presence of noise (just not as efficiently).
- Predictive coding can be used in both lossless and lossy compression schemes.

Predictive coding

- Encoder (transmitter – Alice)

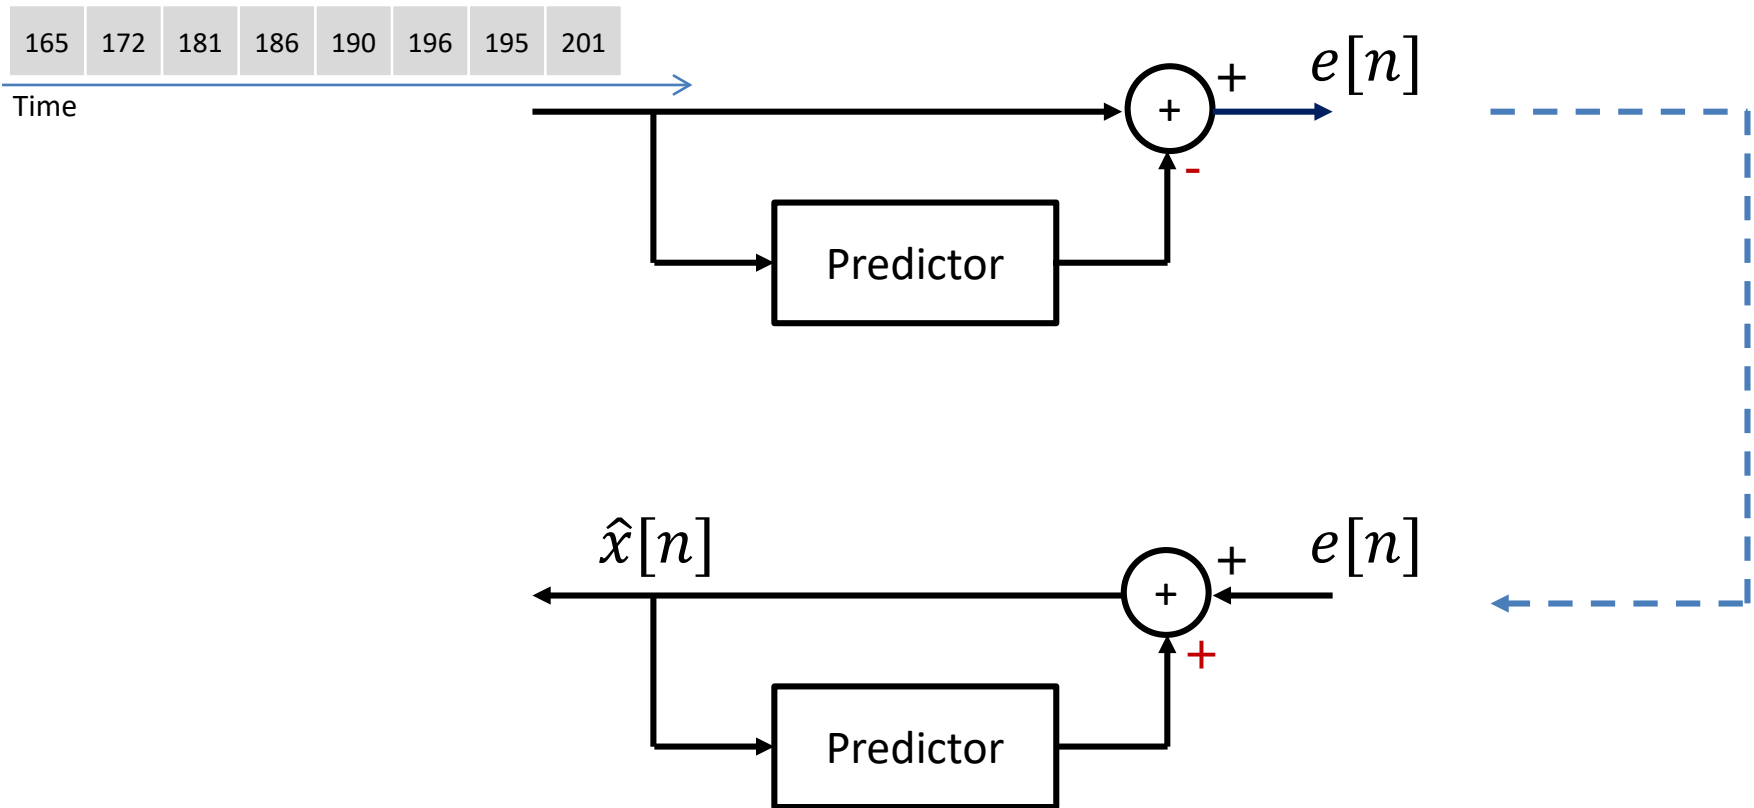


- Decoder (receiver – Bob)



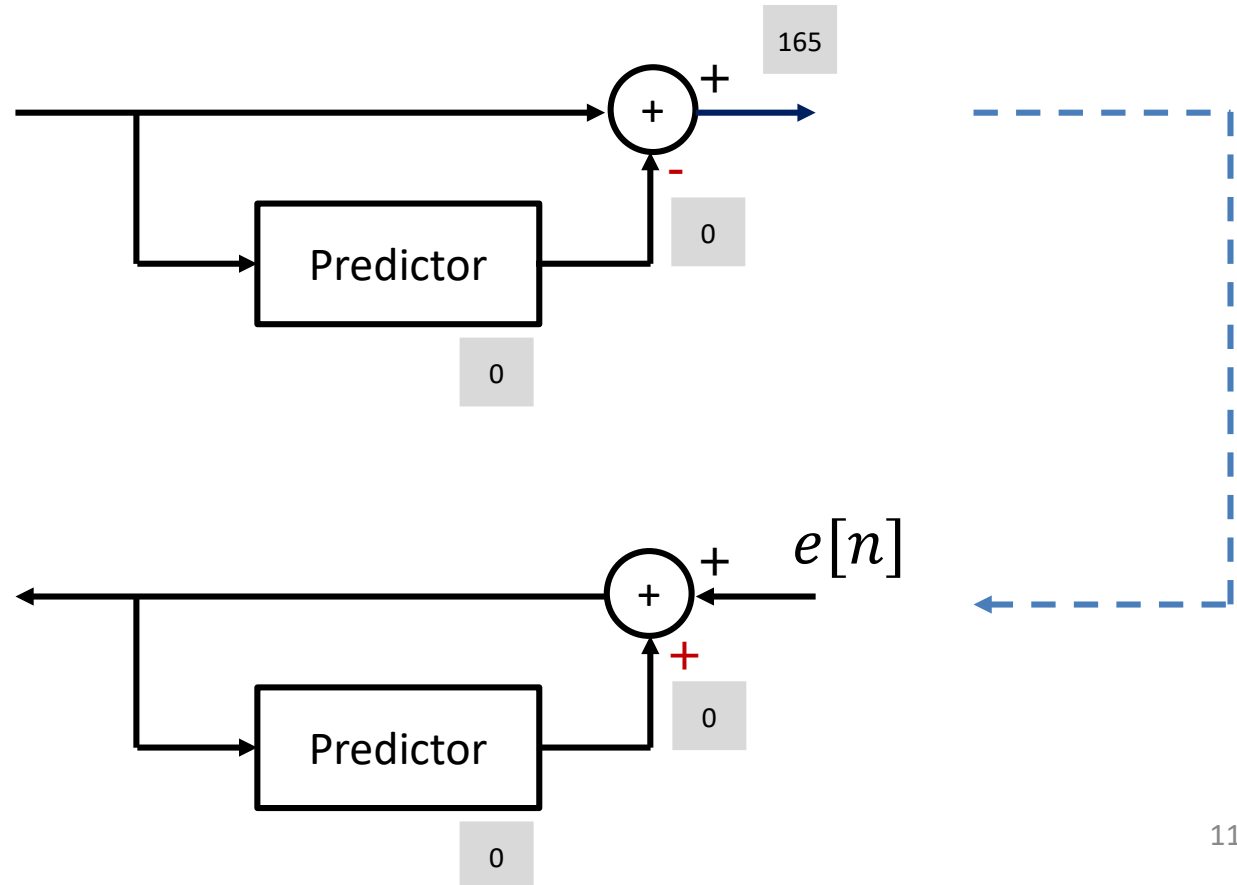
Predictive coding

- Ideal predictor, which **exactly** predict current value based on previous ones



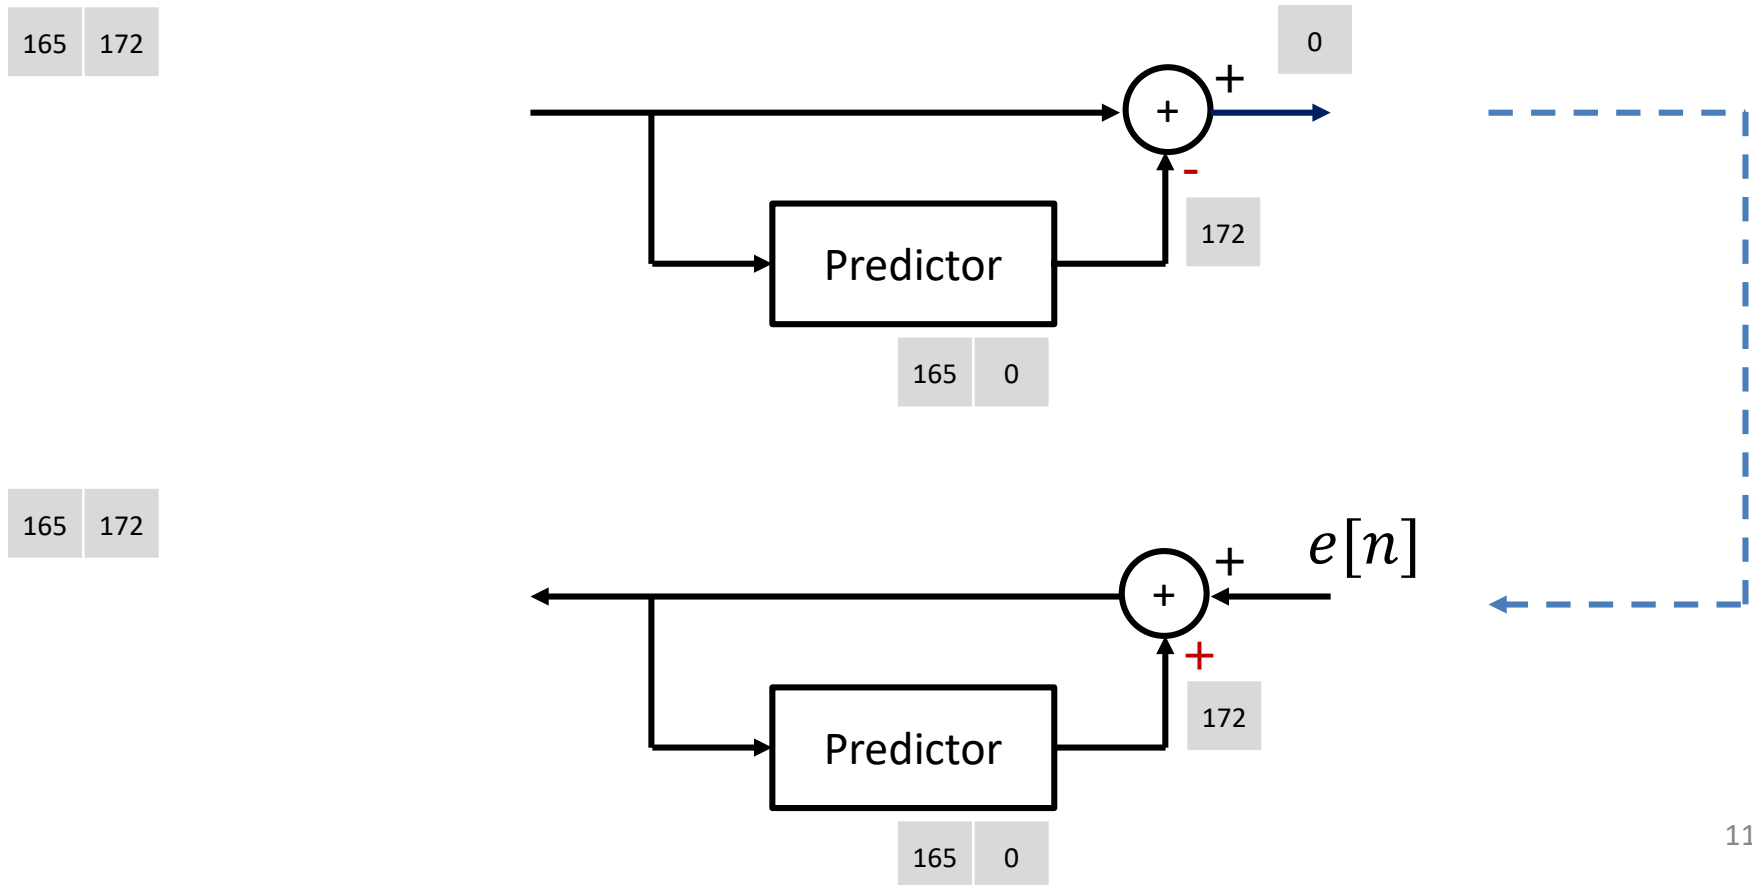
Predictive coding

- Ideal predictor, which **exactly** predict current value based on previous ones



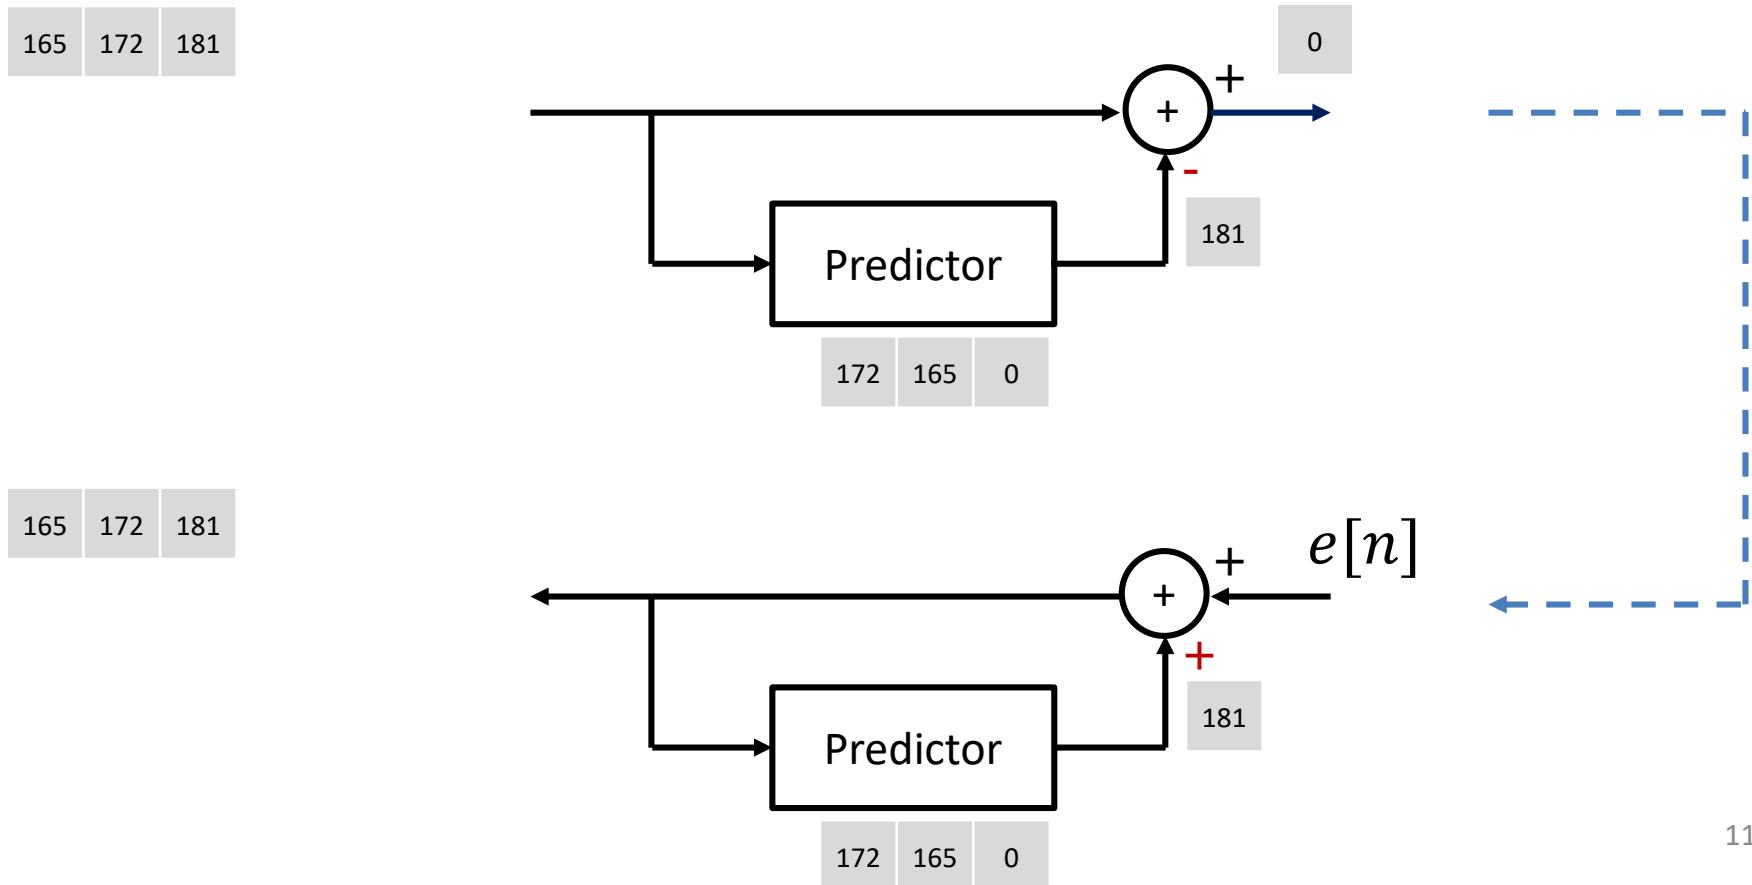
Predictive coding

- Ideal predictor, which **exactly** predict current value based on previous ones



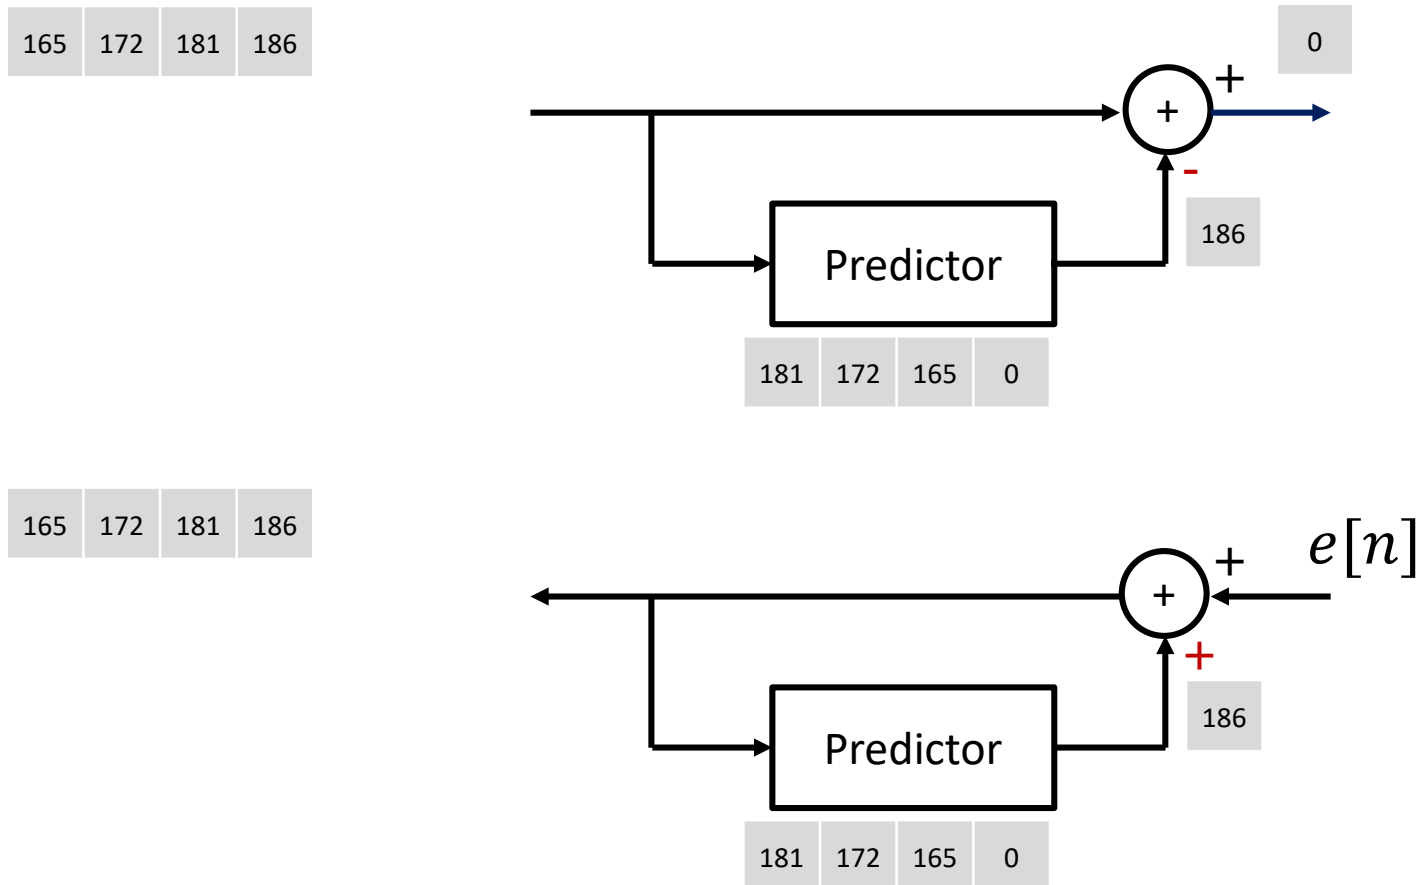
Predictive coding

- Ideal predictor, which **exactly** predict current value based on previous ones



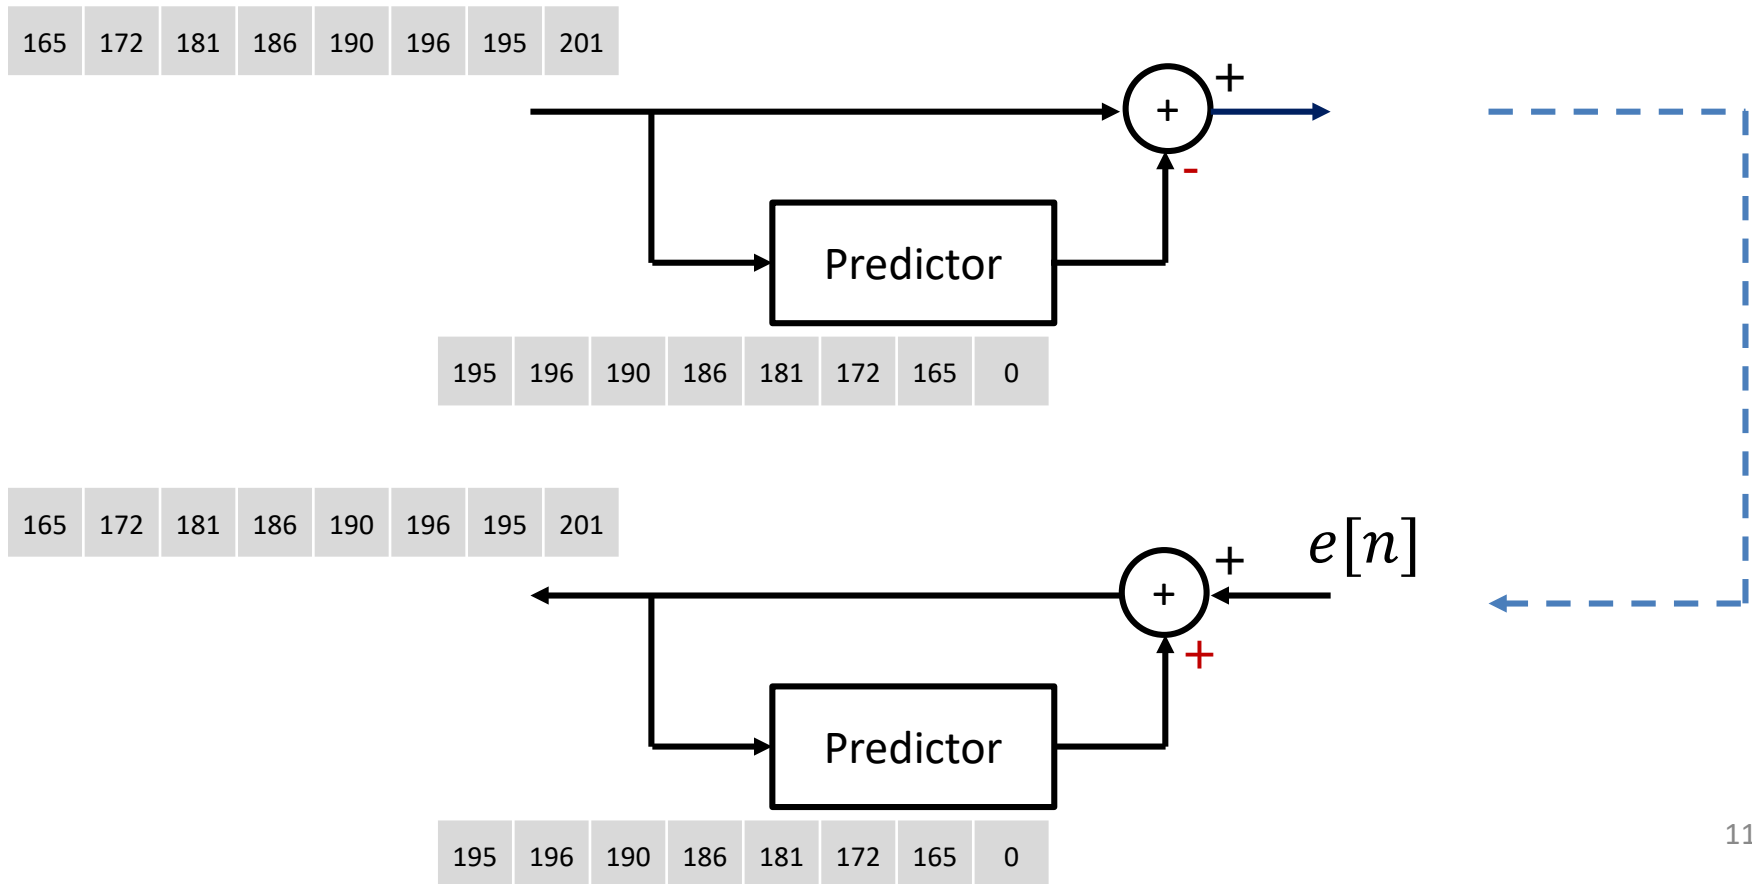
Predictive coding

- Ideal predictor, which **exactly** predict current value based on previous ones



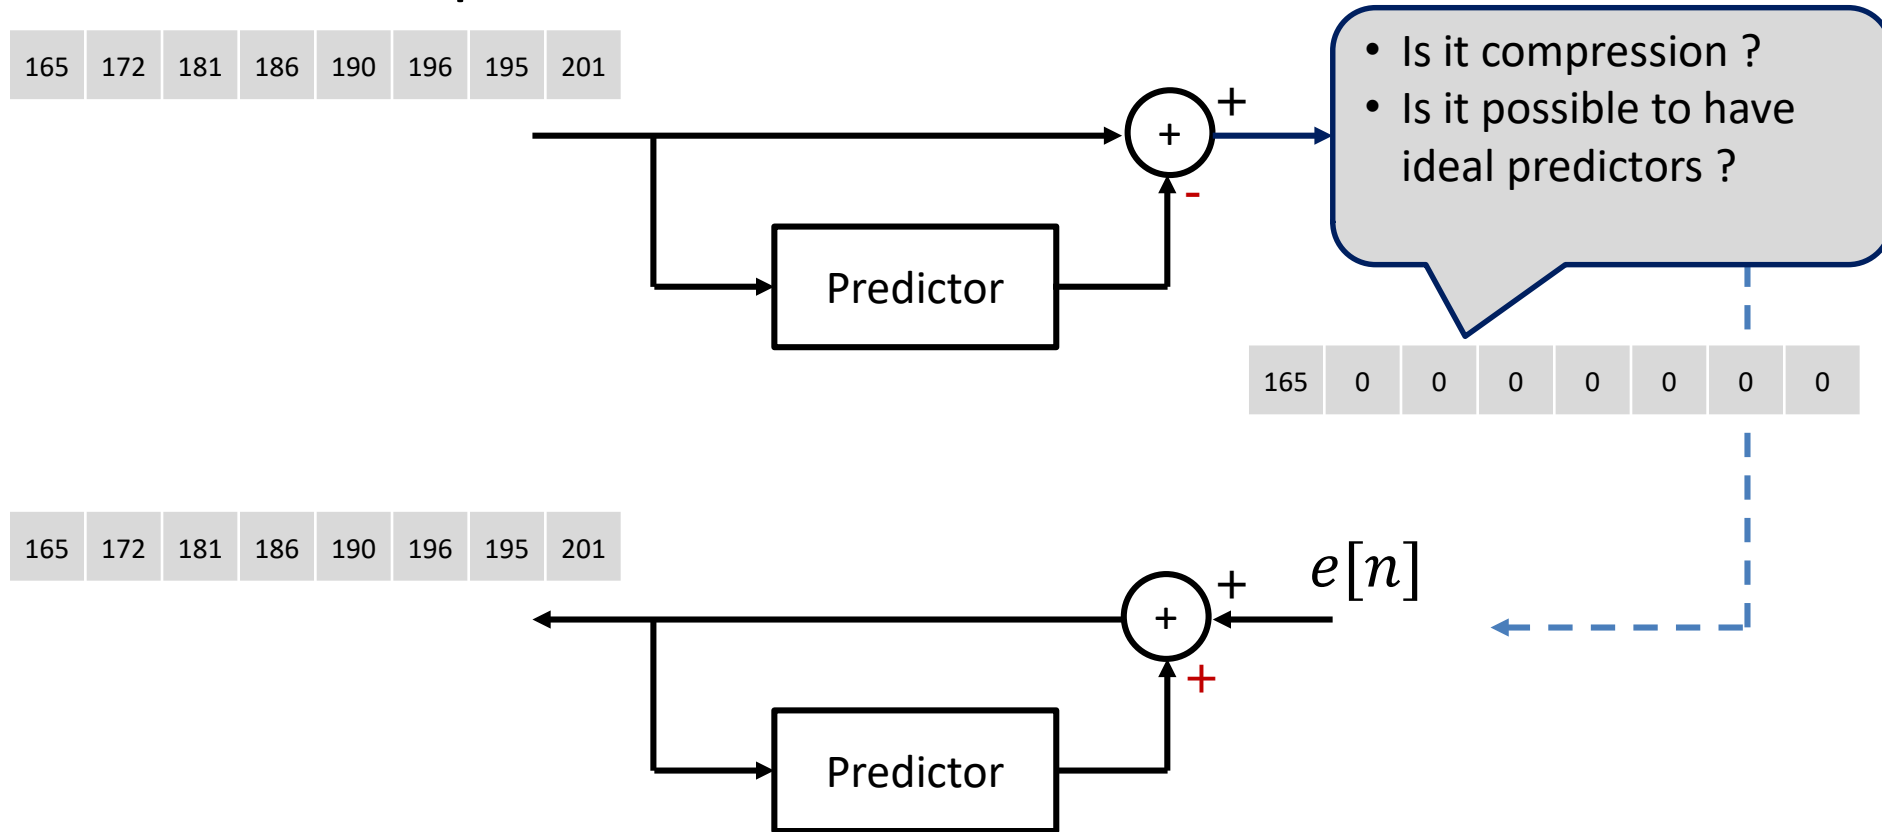
Predictive coding

- Ideal predictor, which **exactly** predict current value based on previous ones



Predictive coding

- Ideal predictor, which **exactly** predict current value based on previous ones




Predictive coding for images


- linear prediction (running on rows)

$$e(r, c) = x(r, c) - x(r, c - 1)$$

Original image entropy = 8.00



DPCM image entropy = 0.10



Predictive coding for images

- Raster-scan DPCM coding (differential pulse code modulation)

165	172	181	186	190	196	195	201
169	176	184	187	192	193	194	195
169	173	182	187	190	193	189	190
173	177	182	185	191	189	189	188
168	173	179	182	189	187	188	190
169	170	175	180	183	184	185	189
166	169	173	176	181	180	186	184
171	168	167	176	176	180	177	181

$$p(r, c) = f(x(r, c-1), x(r-1, c-1), x(r-1, c))$$

$$e(r, c) = x(r, c) - p(r, c)$$

Predictive coding for images

- Raster-scan DPCM coding

165	172	181	186	190	196	195	201
169	176	184	187	192	193	194	195
169	173	182	187	190	193	189	190
173	177	182	185	191	189	189	188
168	173	179	182	189	187	188	190
169	170	175	180	183	184	185	189
166	169	173	176	181	180	186	184
171	168	167	176	176	180	177	181

$$p(r, c) = f(x(r, c-1), x(r-1, c-1), x(r-1, c))$$
$$e(r, c) = x(r, c) - p(r, c)$$

Predictive coding for images

- Raster-scan DPCM coding

$$e(r, c) = x(r, c) - \frac{x(r, c-1) + x(r-1, c-1) + x(r-1, c)}{3}$$

DPCM output



Original



Predictive coding for images

- Raster-scan DPCM coding

$$e(r, c) = x(r, c) - \frac{x(r, c-1) + x(r-1, c-1) + x(r-1, c)}{3}$$



Psychovisual redundancy

Psychovisual redundancy

- Human **perception** of the information of an **image** does not involve quantitative **analysis** of **every pixel**
- **Pixel values** can be **modified** up to a given extent without **significant** subjective **degradation**
- **Proper** alteration should involve **psychovisually redundant** information
 - The image is **irreversibly altered**

Psychovisual redundancy



- Psychovisual redundancy **stem** from that the **human eye** does **not respond equally** to all visual information.
- An **observer searches** for **distinct features** and **mentally combines** them into recognizable **objects**.
 - Use of **prior knowledge** for **interpretation** (face, wall, ...)
- In **this process certain** information is relatively **less important** than other — this information is called psychovisually redundant.
- If the wall were slightly different this could not be perceived

LOSSY DATA COMPRESSION (Non reversible algorithms)



Introduction

- Advantages of lossy compression:
 - Higher compression ratios
- Disadvantages:
 - The decoded signal is not exactly conform to the original

Examples of lossy compression

- Speech/Audio:
 - GSM/UMTS/WCDMA (multi-rate) speech compression
 - MP3 audio
- Image compression:
 - JPEG
 - JPEG2000
- Video compression
 - H.264/AVC , H.264/SVC
 - H.265 (HEVC)


Lossy compression

- Types of lossy compression
 - Predictive coding
 - Transform coding
 - Model-based coding
 - ...
- Lossy quality measures:
 - PSNR
 - SNR
 - Subjective measures

transform coding

General transform coding scheme

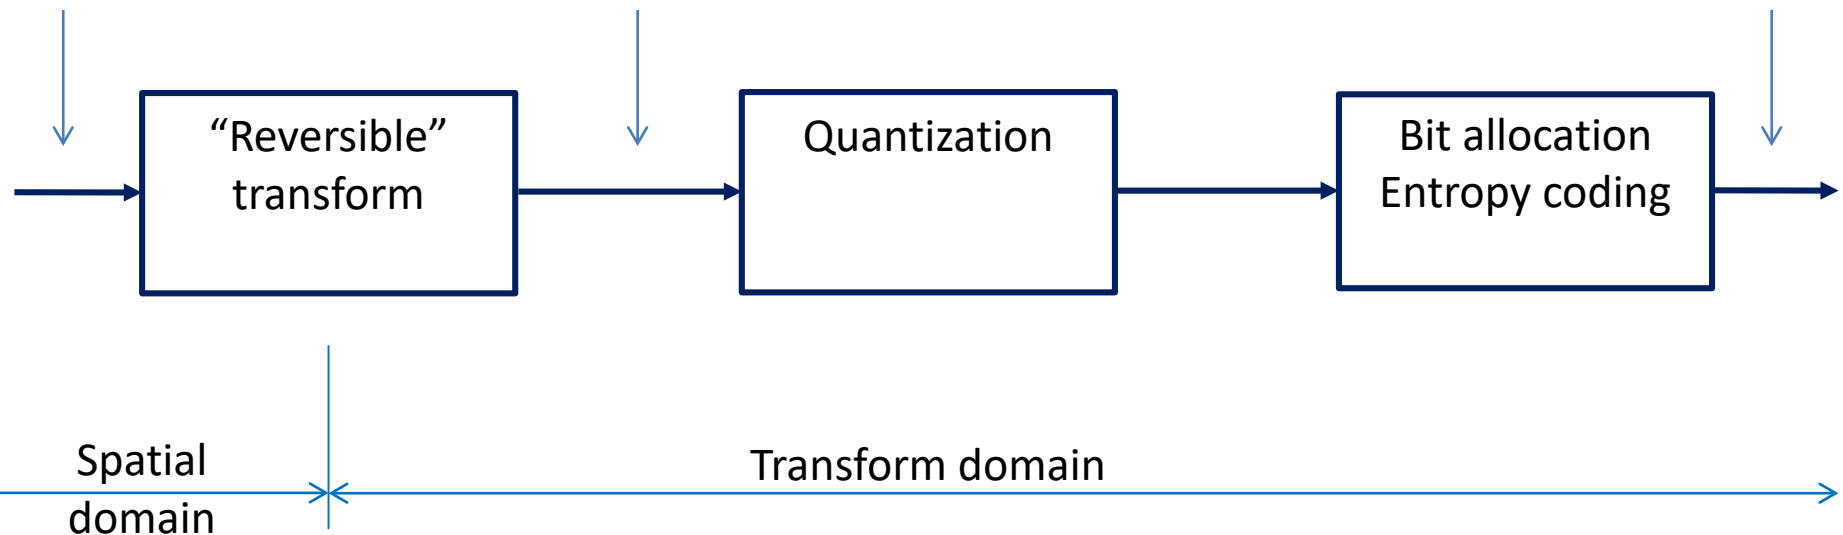
Original



The transformation aims to reduce the inter-pixel correlation

- KLT (Karhunen-Loeve)
- DCT (Discrete cosine)
- DWT (Discrete wavelet)

165	172	181	186	190	196	195	201
-----	-----	-----	-----	-----	-----	-----	-----



Example of 2-D DCT

Image

DCT

139	144	149	153	155	155	155	155
144	151	153	156	159	156	156	156
150	155	160	163	158	156	156	156
159	160	162	160	160	159	159	159
159	160	161	162	162	155	155	155
161	161	161	161	160	157	157	157
162	162	161	163	162	157	157	157
162	162	161	162	163	158	158	158

236	-1	-12	-5	2	-2	-3	1
-23	-17	-6	-3	-3	0	0	-1
-11	-9	-2	2	0	-1	-1	0
-7	-2	0	1	1	0	0	0
-1	-1	1	2	0	-1	1	-1
2	0	2	0	-1	1	1	-1
-1	0	0	-1	0	2	1	-1
-3	2	-4	-2	2	1	-1	0

2-D DCT



- The two-dimensional DCT (2D-DCT) could be **written** as :

$$F(k, l) = \alpha(k)\alpha(l) \sum_{n=0}^{N-1} \sum_{m=0}^{N-1} x(n, m) \cos\left[\frac{(2n+1)k\pi}{2N}\right] \cos\left[\frac{(2m+1)l\pi}{2N}\right]$$

$k, l = 0, \dots, N-1$

$$\alpha(0) = \sqrt{\frac{1}{N}} \quad \alpha(k) = \sqrt{\frac{2}{N}}$$

Why the need for transform coding ?



code

Decorrelation property of the DCT

- It can be noticed that the **coefficients** in the **transformed domain** are much more **unbalanced**
- Or equivalently we could say in the transformed domain few coefficients concentrate most of the signal energy → therefore can be more **efficiently compressed**
- This property holds for most frequency transforms, if the image has low pass characteristics (smooth areas)

Joint Photographic Expert Group -JPEG-

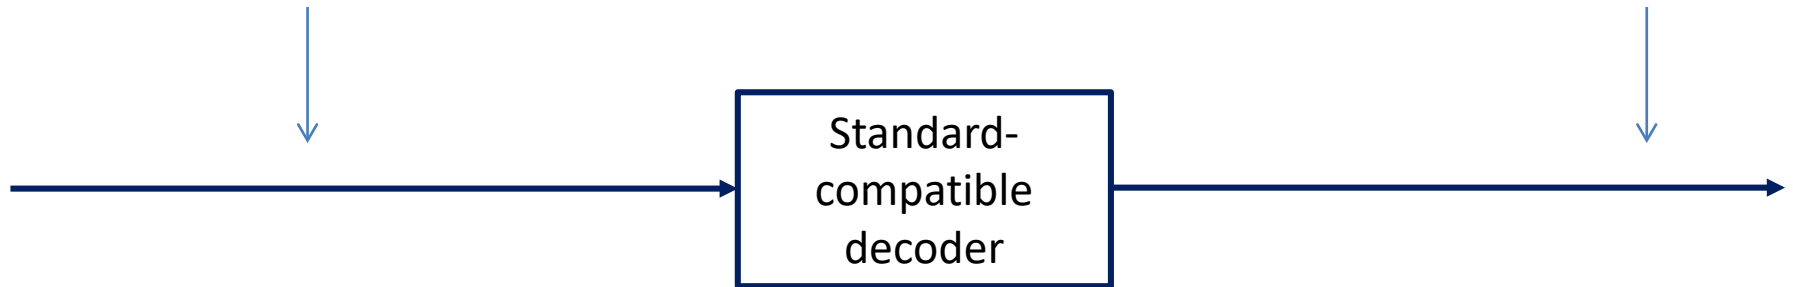
The JPEG standard

- JPEG stands for “Joint Photographic Expert Group”.
- Voted as international standard in 1992
- This standardized image compression scheme is designed to work on full-color or gray-scale
- It specifies the decoder and the codestream syntax

The JPEG standard

- Most image standard specifies the decoder and the codestream syntax

165	172	181	186	190	196	195	201
-----	-----	-----	-----	-----	-----	-----	-----



It is **not** very important **how** you **compress** (implement the standard), what matter is to generate a standard-compatible stream

JPEG coding steps

- Transformation of the image into a suitable color space
- Application of a 8x8 blocks DCT
- Quantization
- *zig-zag* reading
- Entropy (lossless) coding

JPEG: 8x8 DCT

- 8x8 DCT transform

image size is 256 x 256




- The image is processed as 8x8 blocks

JPEG: 8x8 DCT

- 8x8 DCT transform


image size is 256 x 256



- Each 8x8 block of pixels is separately DCT transformed

JPEG: 8x8 DCT


- One block pixel



96	99	99	95	86	95	107	121
97	96	93	107	99	98	109	124
98	96	95	98	98	95	109	118
98	99	98	97	96	95	108	125
98	97	97	95	97	97	106	121
99	100	97	99	95	101	111	125
98	104	98	99	97	99	110	125
103	103	95	96	92	107	111	128

JPEG: 8x8 DCT

- 8x8 DCT transform



818.0	-46.6	44.4	-26.7	12.0	-3.0	-1.9	-6.0
-9.1	1.7	-5.1	-4.0	3.4	4.5	-2.0	2.2
2.8	-0.7	4.2	-0.3	-3.9	1.3	1.2	-6.7
-2.5	1.2	0.7	-3.7	-3.1	0.2	1.4	-0.2
-3.8	1.7	6.4	-1.0	-4.3	-1.7	4.9	0.2
-1.9	3.0	2.7	-4.5	-3.9	-1.4	-0.9	2.7
-2.5	0.7	3.4	2.3	-4.1	1.8	2.1	0.5
-4.9	3.0	-1.5	1.5	-1.6	-0.8	-0.2	3.4

- The **DCT coefficients** are real valued, so forward and inverse DCT are limited by computer precision
- No truly **lossless** compression is possible if the **non-integer DCT** is used


JPEG: 8x8 DCT

- Lenna in the 8x8 DCT domain

image size is 256 x 256



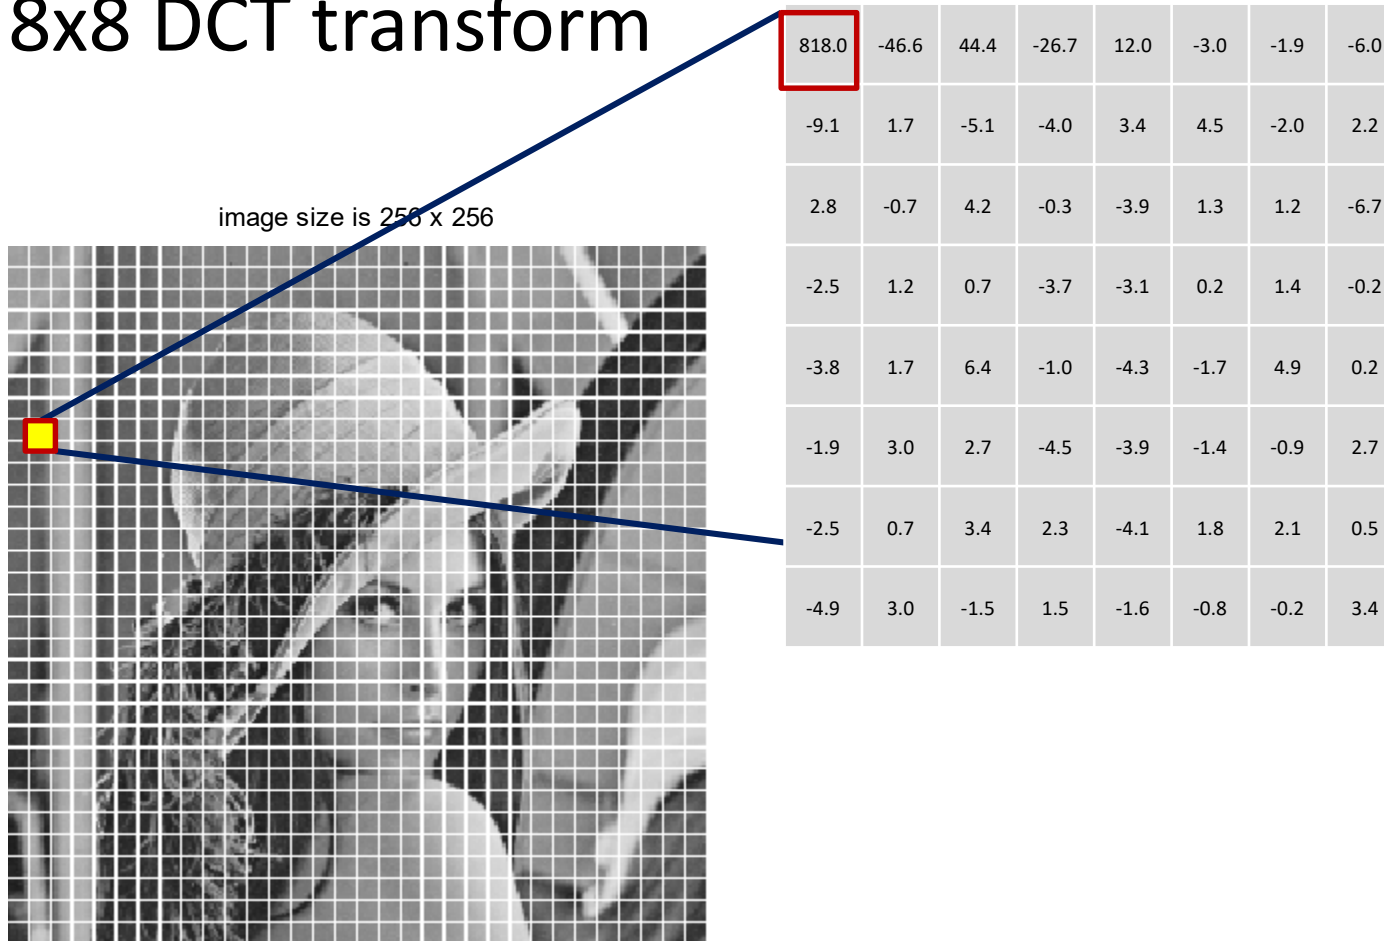
the 8 x 8 2D DCT of the image



Why ?

JPEG: 8x8 DCT

- 8x8 DCT transform




JPEG: 8x8 DCT

- 8x8 DCT transform


This format is not used in JPEG, it is just illustrative, however, it shows that the **DC values** are still **correlated**

the 8 x 8 2D DCT of the image (Natural order)



Natural order

the 8 x 8 2D DCT of the image (Reordered)



Reordered (DC coeff. grouped)

JPEG: quantization

- This is the truly lossy stage of JPEG
- The *quantization matrix* is an 8 by 8 matrix of *step sizes*, with one element for each DCT coefficient
- *Step sizes* will be small for *low frequencies*, and large for *high frequencies*.
- The *quantizer divides* the DCT coefficient by its *corresponding* quantization *step*, then *rounds* to the *nearest integer*

Quantization example

235.6	-1.0	12.1	-5.2	2.1	-1.7	-2.7	1.3
-22.6	-17.5	-6.2	-3.2	-2.9	-0.1	0.4	-1.2
-10.9	-9.3	-1.6	1.5	0.2	-0.9	-0.6	-0.1
-7.1	-1.9	0.2	1.5	0.9	-0.1	0.0	0.3
-0.6	-0.8	1.5	1.6	-0.1	-0.7	0.6	1.3
1.8	-0.2	1.6	-0.3	-0.8	1.5	1.0	-1.0
-1.3	-0.4	-0.3	-1.5	-0.5	1.7	1.1	-0.8
-2.6	1.6	-3.8	-1.8	1.9	1.2	-0.6	-0.4

Original 8x8 block

Quantization table

$Q =$

16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	56	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99

Quantization example

[illegible]

Quantized 8x8 block


Dequantized 8x8 block

[illegible]

JPEG: quantization

- 8x8 DCT transform

image size is 256 x 256



818.0	-46.6	44.4	-26.7	12.0	-3.0	-1.9	-6.0
-9.1	1.7	-5.1	-4.0	3.4	4.5	-2.0	2.2
2.8	-0.7	4.2	-0.3	-3.9	1.3	1.2	-6.7
-2.5	1.2	0.7	-3.7	-3.1	0.2	1.4	-0.2
-3.8	1.7	6.4	-1.0	-4.3	-1.7	4.9	0.2
-1.9	3.0	2.7	-4.5	-3.9	-1.4	-0.9	2.7
-2.5	0.7	3.4	2.3	-4.1	1.8	2.1	0.5
-4.9	3.0	-1.5	1.5	-1.6	-0.8	-0.2	3.4

JPEG: quantization






image size is 256 x 256

2D-DCT

```

» im = double(imread('lenna256.bmp'));
» [r,c] = size(im);
» mask = zeros(r,c);
» for ri = 64 + [1 : 8] -1
»     for ci = 8+[1:8]-1
»         mask(ri,ci) = 1;
»     end;
» end;
» blk = reshape(im(mask == 1), 8,8);
» subplot(1,3,1); image(blk); colormap(gray(255)); axis off;
» blkdct = dct2(blk);
» Q_step = 16;
» Qblkdct = round(blkdct/Q_step);
» subplot(1,3,2); image(255*(Qblkdct ~= 0));
» colormap(gray(255)); axis off; title(sprintf('Quantization step
» is %d', Q_step));
» iblkdct = idct2(Q_step * Qblkdct);
» subplot(1,3,3); image(iblkdct); colormap(gray(255)); axis off;
» truesize([140 140])
    
```

JPEG: quantization



Q-step=4





De-quntized

```

» im = double(imread('lenna256.bmp'));
» [r,c] = size(im);
» mask = zeros(r,c);
» for ri = 64 + [1 : 8] -1
»     for ci = 8+[1:8]-1
»         mask(ri,ci) = 1;
»     end;
» end;
» blk = reshape(im(mask == 1), 8,8);
» subplot(1,3,1); image(blk); colormap(gray(255)); axis off;
» blkdct = dct2(blk);
» Q_step = 4;
» Qblkdct = round(blkdct/Q_step);
» subplot(1,3,2); image(255*(Qblkdct ~= 0));
» colormap(gray(255)); axis off; title(sprintf('Quantization step
» is %d', Q_step));
» iblkdct = idct2(Q_step * Qblkdct);
» subplot(1,3,3); image(iblkdct); colormap(gray(255)); axis off;
» truesize([140 140])
    
```

JPEG: quantization



Q-step=16



De-quntized

```

» im = double(imread('lenna256.bmp'));
» [r,c] = size(im);
» mask = zeros(r,c);
» for ri = 64 + [1 : 8] -1
»     for ci = 8+[1:8]-1
»         mask(ri,ci) = 1;
»     end;
» end;
» blk = reshape(im(mask == 1), 8,8);
» subplot(1,3,1); image(blk); colormap(gray(255)); axis off;
» blkdct = dct2(blk);
» Q_step = 16;
» Qblkdct = round(blkdct/Q_step);
» subplot(1,3,2); image(255*(Qblkdct ~= 0));
» colormap(gray(255)); axis off; title(sprintf('Quantization step
» is %d', Q_step));
» iblkdct = idct2(Q_step * Qblkdct);
» subplot(1,3,3); image(iblkdct); colormap(gray(255)); axis off;
» truesize([140 140])
    
```

JPEG: quantization

- JPEG quantization matrix (table)

$$Q = \begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 56 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$

JPEG: quantization effects



- Large quantization **step** drive small **coefficients** down to **zero**
- The result: many **high frequency** coefficients become **zero**, and therefore easier to code
- The **low frequency** coefficients undergo only **minor adjustment**
- The quantization matrix can be rescaled by multiplication by a **quality factor(QP)**

Quality factor

- When using **JPEG images**, one can **set** the **quality** of the image from very **low quality** to very **high quality**
- The file size varies inversely with the quality of the image.
- There are basically four "standard" levels of JPEG: Low, med, high, max


JPEG encoding

- Quantized DC values are coded by DPCM from block to block to remove the residual correlation.



JPEG encoding

- AC coefficients:



- Zig-zag reordering is performed to achieve large runs of zeros
 - Encoding of zero-runs
 - Entropy coding (Huffman)

JPEG encoding

- DC values:
 - Quantized DC values are coded by DPCM from block to block to remove the residual correlation.
- AC coefficients:
 - Zig-zag reordering is performed to achieve large runs of zeros
 - Encoding of zero-runs
 - Entropy coding (Huffman)

JPEG decoding

- The encoding steps are reversed
- Huffman decoding
- Run length decoding
- Coefficient de-quantization (each coefficient multiplied by the quantum)
- Inverse DCT

JPEG performance



A photo of a cat with the compression rate decreasing, and hence quality increasing, from left to right

From Wikipedia.org

JPEG performance



Quality max - Size: 61k




Quality med - Size: 14k



Quality low - Size: 4k

JPEG Baseline Example




Reconstructed
images

Differences between
original images and
reconstructed images
(RMS errors 2.5 on
left, 4.4 on right)

Details of
reconstructed images

Left: JPEG Compression with, CR = 18:1
Right: JPEG Compression with, CR = 42:1

JPEG 2000 Example



Reconstructed
images

Differences between
original images and
reconstructed images
(RMS errors 3.7 on left,
5.9 on right)

Details of
reconstructed images

Left: JPEG 2000 compression, CR = 42:1

Right: JPEG 2000 compression, CR = 88:1

JPEG performance



Original image

Encoded @ 24 bits per pixel

JPEG performance



Quality 95/100
3.926 bits per pixel (bpp)
 $CR = 24/3.926 = 6.1$

JPEG performance



Quality 50/100

1.067 bits per pixel (bpp)

CR = 22.5

JPEG performance



Quality 25/100

0.705 bits per pixel (bpp)

CR = 34.0

JPEG performance



Quality 5/100 (min.useful)

0.291 bits per pixel (bpp)

CR = 82.5

Thanks